

Mielage Prediction Project

```
In [2]: #import library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: #import data
df=pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	usa	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	usa	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	usa	ford torino

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64  
2   displacement     398 non-null   float64
3   horsepower       392 non-null   float64
4   weight           398 non-null   int64  
5   acceleration     398 non-null   float64
6   model_year       398 non-null   int64  
7   origin           398 non-null   object  
8   name             398 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
In [6]: df.nunique()
```

```
Out[6]: mpg          129
cylinders          5
```

```
displacement    82
horsepower       93
weight          351
acceleration     95
model_year       13
origin           3
name            305
dtype: int64
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
count	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	76.010050
std	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689	3.697627
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	73.000000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000	79.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000

```
In [8]: df.corr()
```

```
Out[8]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
mpg	1.000000	-0.775396	-0.804203	-0.778427	-0.831741	0.420289	0.579267
cylinders	-0.775396	1.000000	0.950721	0.842983	0.896017	-0.505419	-0.348746
displacement	-0.804203	0.950721	1.000000	0.897257	0.932824	-0.543684	-0.370164
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361
weight	-0.831741	0.896017	0.932824	0.864538	1.000000	-0.417457	-0.306564
acceleration	0.420289	-0.505419	-0.543684	-0.689196	-0.417457	1.000000	0.288137
model_year	0.579267	-0.348746	-0.370164	-0.416361	-0.306564	0.288137	1.000000


```
In [12]: #remove missing value
df=df.dropna()
df.columns
```

```
Out[12]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
               'acceleration', 'model_year', 'origin', 'name'],
              dtype='object')
```


```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 392 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   mpg              392 non-null   float64
1   cylinders        392 non-null   int64  
2   displacement     392 non-null   float64
3   horsepower       392 non-null   float64
4   weight           392 non-null   int64  
5   acceleration     392 non-null   float64
6   model_year       392 non-null   int64  
7   origin           392 non-null   object  
8   name             392 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```

```
In [14]: #Data visualization
sns.pairplot(df,x_vars=['mpg','displacement','horsepower','weight',
                        'acceleration'],y_vars=['mpg']);
```



```
In [15]: sns.regplot(x='displacement',y='mpg',data=df);
```



```
In [16]: #define x and y
df.columns
```

```
Out[16]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
               'acceleration', 'model_year', 'origin', 'name'],
              dtype='object')
```

```
In [17]: y=df['mpg']
x=df[['displacement', 'horsepower', 'weight',
      'acceleration']]
```

```
In [18]: #split
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=2529)
```

```
In [19]: #model
from sklearn.linear_model import LinearRegression
model=LinearRegression()
```

```
In [20]: model.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

```
In [21]: model.intercept_
```

```
Out[21]: 45.84620249789292
```

```
In [22]: model.coef_
```

```
Out[22]: array([-0.007859 , -0.05202824, -0.0048651 , -0.05998945])
```

```
In [23]: #prediction  
y_pred=model.predict(x_test)  
y_pred
```

```
Out[23]: array([18.45029029, 15.11872575, 14.25951901, 23.63777162, 29.77227939,  
23.78289678, 26.46274613, 24.63477759, 15.10361067, 11.92089347,  
24.03667612, 28.03774179, 31.7791986 , 31.04942136, 18.34939414,  
19.34562679, 28.14901371, 32.26833498, 31.23336778, 27.1706607 ,  
18.90264044, 22.69158865, 26.30616149, 32.53334114, 20.7455229 ,  
8.43604922, 21.96939005, 18.16644283, 24.9187207 , 14.95041612,  
23.27573018, 17.10008397, 9.28416594, 30.02859334, 20.49341373,  
29.16402497, 24.1851619 , 21.82468561, 10.45764414, 12.99758931,  
21.55287965, 19.9763373 , 5.81701795, 17.83479167, 22.69872144,  
29.39987303, 13.2638446 , 25.84303202, 29.29886179, 22.44116443,  
22.30857618, 16.57432268, 24.06827363, 30.19019859, 10.04817173,  
9.3533171 , 28.14495274, 23.67665202, 20.07936568, 30.77322956,  
20.95405256, 26.72684739, 23.16157669, 14.10789682, 24.37223149,  
26.84731155, 15.26437637, 24.21355 , 30.81705563, 14.86794633,  
27.5428809 , 24.35148953, 10.75013125, 30.29658039, 30.95694009,  
27.35893598, 31.26808388, 10.29239165, 27.64504505, 16.41746006,  
25.5910977 , 29.48584659, 14.83958315, 32.76319208, 30.34965318,  
30.95305498, 14.61576534, 27.04413659, 26.74989971, 29.0983602 ,  
32.55952574, 29.50578249, 31.70671628, 31.69454341, 21.58369883,  
31.71427871, 26.19466037, 28.94617784])
```

```
In [25]: #accuracy  
from sklearn.metrics import mean_absolute_percentage_error  
mean_absolute_percentage_error(y_test,y_pred)
```

```
Out[25]: 0.14486145216628077
```

```
In [ ]:
```